Mumtozbek Akhmadjonov

oltinsahifa@gmail.com | akhmadzhonov.mk@phystech.edu | +7 (925) 662-52-01
github.com/mumtozee | linkedin.com/in/imumtozee

Education

Moscow Institute of Physics and Technology

M.S. in Computer Science and Data Science — *Sep. 2023 - Aug. 2025*

Relevant Courses: Computer Vision, Natural Language Processing, Reinforcement Learning

GPA: 5.0/5.0, Top-1% of the class

B.S. in Computer Science and Applied Mathematics — *Sep. 2019 - Aug. 2023*

Relevant Courses: Machine Learning, Linear Algebra, Calculus I and II, Convex Optimizations, Probability Theory and Mathematical Statistics, Random Processes

Experience

Yandex

ML Engineer — *Jan. 2024 - Jul. 2025*

- Increased the pretraining dataset quality (+4% ANA, +8% CapScore, +2% CLIP-Score) of YandexVLM by creating and implementing CapMerge: a method to enhance image captions by filtering out visual hallucinations and increasing granularity.
- Increased model performance on popuplar VLM benchmarks (RealWorldQA, HallB, MMMU etc.) by pretraining on enhanced image captions. (+1.8% on average)
- Reproduced ShareGPT4V VLM and further improved it's captioning abilities.
- Made YaART text-to-image model better by improving the pretraining dataset.

Intern ML Engineer — *Sep. 2023 - Dec. 2023*

- Finetuned Llama 2 in a multimodal setting using speech tokens from both CommonVoice and CoVoST2 datasets to make ASR and AST. Achieved a BLEU score similar to Google's AudioPaLM-1 8B on CoVoST2.
- Introduced GPTQ, a SotA int4 quantization algorithm, for the Seq2Seq ASR model from the production increasing it's quality (WERp -2.87% on hard samples) and inference speed. Reduced the required GPU memory by 3 times.

DeepPavlov

Junior DL Researcher — *Sep. 2022 - Jul. 2023*

- Developed approaches to evaluate and enhance Dialogue Graph Auto-Construction (DGAC) method on different dialogue datasets. Co-authored a paper accepted to AINL 2023.
- Outperformed the approach described in the paper on user attribute extraction and inference from dialogues with SOTA transformer architectures.

Research Intern — *Feb. 2022 - May. 2022*

- Achieved a similar performance using DialoGPT with special tokens to the approach leveraging projected attention layers to control different dialogue attributes. Reviewed dozens of articles on the related topics on ArXiv.

Publications

XXV International Conference on Neuroinformatics

- Co-authored a conference paper about dialogue response selection enhancement using conversational graphs.

Neuroinformatics 2022 NN workshop

- Presented a poster describing controllable DialoGPT.

Projects

User Attribute Extraction

- Extracting structured persona infromation from conversational data.

RLHF on GPT2

- My attempt to align GPT2 using PPO, DPO and SFT on different datasets, and serve the aligned models using FastAPI and gRPC.

PEFT for BERT

- My attempt to fine-tune BERT-like models on MLM task using HuggingFace PEFT.

Skills

Languages:

Python, C, C++, Java

Technologies & Tools:

PyTorch, PyTorch Lightning, HuggingFace Transformers, Accelerate, PEFT, trl, scikit-learn, numpy, pandas matplotlib, Git, Docker, Flask, FastAPI, MapReduce (YTsaurus, Hadoop)

Achievements

OpenDoors International Student's Olympiad 2023

- winner diploma in Math&AI track
- winner diploma in CS&DS track

Skolkovo Hack 2022

- Top-4 with team "ProDaters".

Tinkoff & MSU Math Contest 2021

- winner diploma

Languages

Uzbek: Native	Russian: Bilingual Proficiency	English: C1	Spanish: A2

Job Preferences

Data Science & Machine Learning Engineer/Researcher:

NLP, Computer Vision, Conversational AI, Generative models (GPT family)